

ECE 763-Project 2

April 8, 2018

Contents

1	Preprocessing Dataset	1
2	Babysitting Neural Network	1
2.1	Architectures	1
2.1.1	Multi-Layer Perceptron	1
2.1.2	Lenet	4
2.1.3	Own Architecture	7
3	Conclusion	9
4	GitHub Repo for Code	9

1 Preprocessing Dataset

The dataset used was a combination of the AFLW, FDDB and LFW datasets. The total number of training images was 10,000 and the number of testing images was 1000, per class. Since this is a binary classification task, the total size of the training set was 20,000 and the total size of the testing set was 2000.

The dimensions of the cropped faces and non-faces were taken to be 64*64.

The glob library was used for easy reading of files. The process is slow on the first run as it builds a list of paths in the RAM, but for subsequent runs, it is extremely fast.

2 Babysitting Neural Network

This section describes the different architectures that were used for the face and non-face dataset and the accuracy and loss curves for them.

A short section describing their performance on a few separate test images and a live video stream is also discussed.

Framework: Tensorflow

Build: From Source with GPU support via CUDA 9.1 and CuDNN 7.1

Platform:Ubuntu 16.04

2.1 Architectures

The Adam Optimizer was used for all the models along with softmax with cross entropy as the loss due to the one-hot encoding of the labels.

2.1.1 Multi-Layer Perceptron

The first architecture tested was a simple MLP with 2 hidden layers.

The training hyperparameters and the code for the model are shown below.

```

1 # Parameters
2 learning_rate = 0.001
3 training_epochs = 300
4 batch_size = 128
5 display_step = 10
6 logs_path = './logs/basic_net/'
7
8 # Network Parameters
9 n_hidden_1 = 2048 # 1st layer number of neurons
10 n_hidden_2 = 1024 # 2nd layer number of neurons
11 n_input = patch_size*patch_size*3 # data input (img shape: 64*64)
12 n_classes = 2
13
14 # tf Graph input
15 X = tf.placeholder("float", [None, n_input], name="X")
16 Y = tf.placeholder("float", [None, n_classes], name="Y")
17 print(X.name)
18 print(Y.name)
19 # Store layers weight & bias
20 with tf.name_scope('weights'):
21     weights = {
22         'h1': tf.Variable(tf.random_normal([n_input, n_hidden_1])),
23         'h2': tf.Variable(tf.random_normal([n_hidden_1, n_hidden_2])),
24         'out': tf.Variable(tf.random_normal([n_hidden_2, n_classes]))
25     }
26
27 with tf.name_scope('biases'):
28     biases = {
29         'b1': tf.Variable(tf.random_normal([n_hidden_1])),
30         'b2': tf.Variable(tf.random_normal([n_hidden_2])),
31         'out': tf.Variable(tf.random_normal([n_classes]))
32     }
33
34
35 # Create model
36 def multilayer_perceptron(x):
37     # Hidden fully connected layer with 256 neurons
38     layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
39     layer_1 = tf.nn.relu(layer_1)
40     # Hidden fully connected layer with 256 neurons
41     layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])
42     layer_2 = tf.nn.relu(layer_2)
43     # Output fully connected layer with a neuron for each class
44     out_layer = tf.matmul(layer_2, weights['out']) + biases['out']
45     print(out_layer.name)
46     return out_layer

```

The following figures show the loss and accuracy curves for the basic MLP. Tensorboard was used to generate the graphs in realtime and was instrumental in hyperparameter tuning.

We can see from the figure that the accuracy increases with time and the loss also decreases, but the curves are not very smooth and the performance could be better. This is because we feed the image into the MLP as a **64*64*3** long 1-D vector. There is a loss of spatial information.

Also, a higher batch size produces a smoother loss curve.

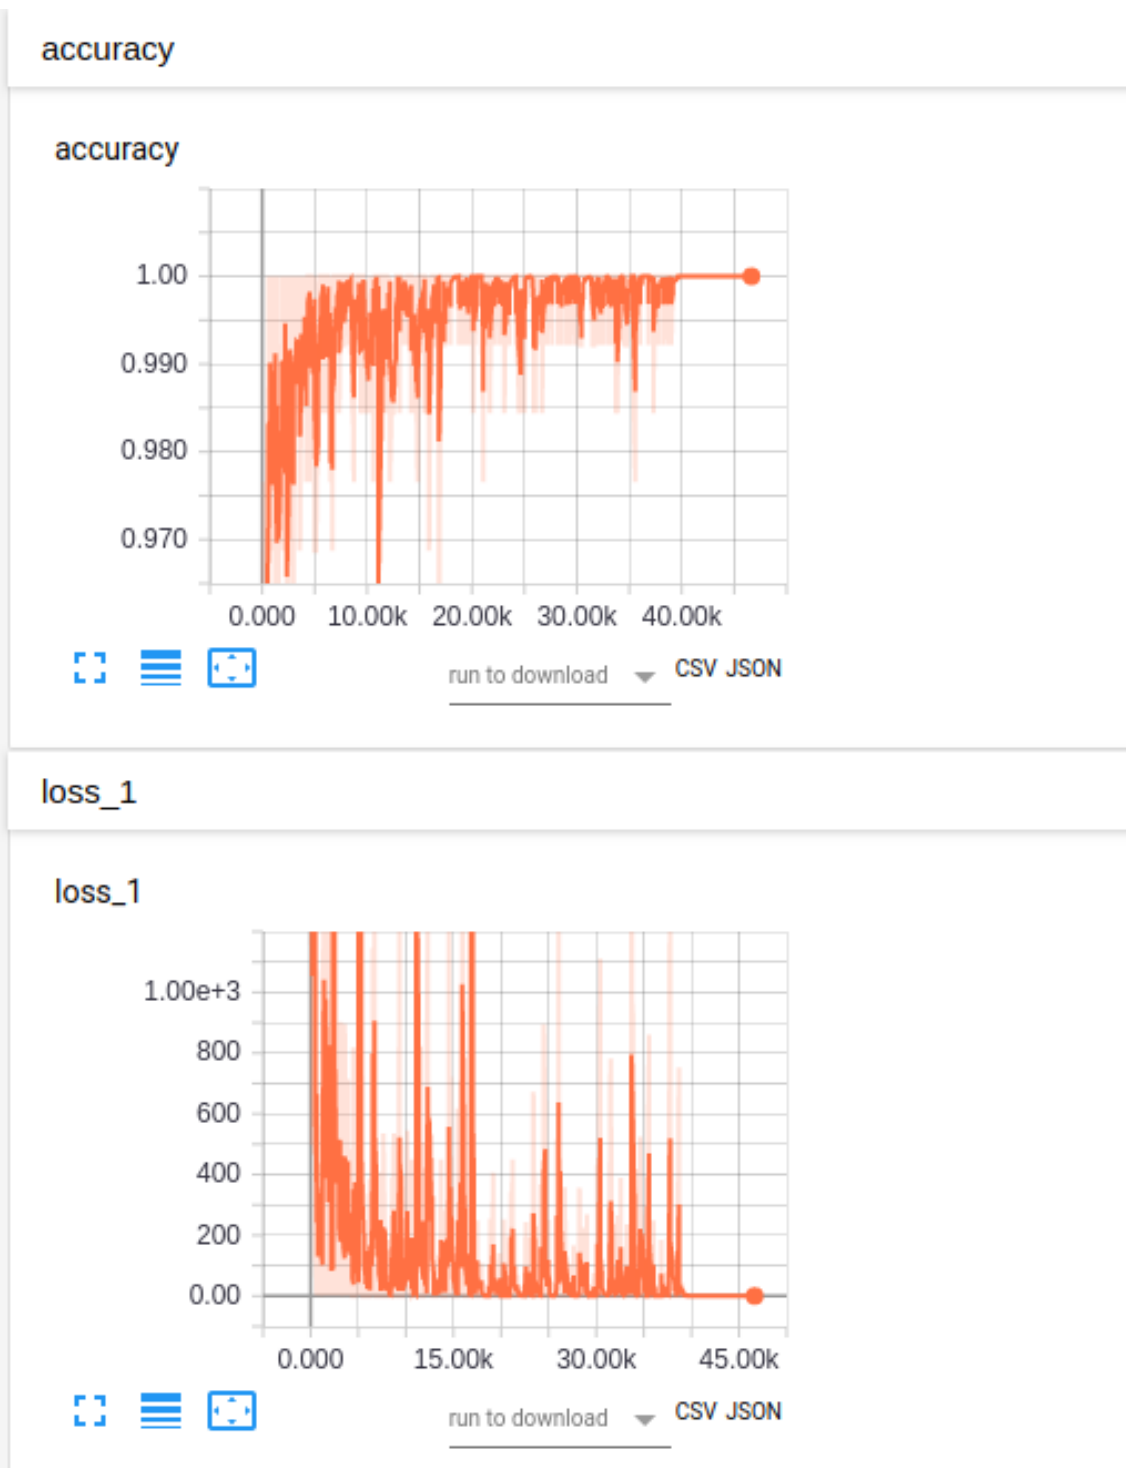


Figure 1: Basic MLP with batch size 128.

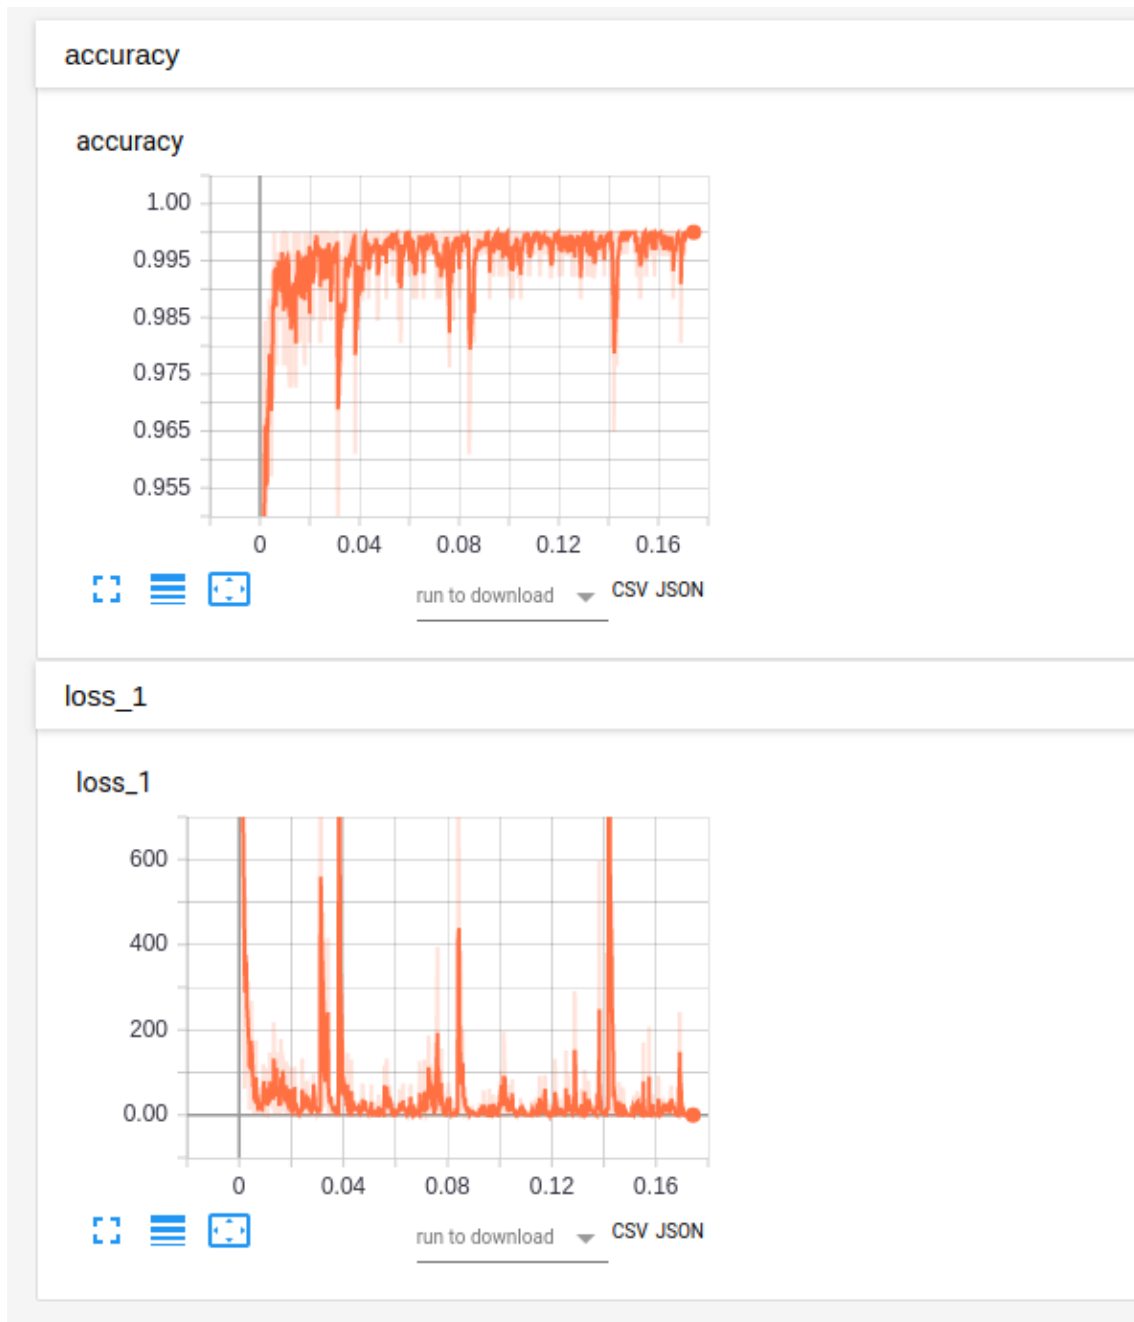


Figure 2: Basic MLP with batch size 256.

Though the MLP is a good start, we can definitely do better. This is why we now see the performance of 2 CNN based approaches. The next architecture is Lenet.

2.1.2 Lenet

The declaration and the hyperparameters are shown in the code section below.

```

1 # Training Parameters
2 learning_rate = 0.01
3 training_epochs = 1000
4 batch_size = 128
5 display_step = 10
6 threshold=0.01

```

```

7 logs_path='./logs/cnn/'
8
9 # Network Parameters
10 num_input = patch_size*patch_size*3 # data input (img shape: patch_size*patch_size)
11 num_classes = 2 # total classes
12 dropout = 0.70 # Dropout, probability to keep units
13
14 # tf Graph input
15 X = tf.placeholder(tf.float32, [None, num_input]) #batch_size, num_input
16 Y = tf.placeholder(tf.float32, [None, num_classes]) #batch_size, num_classes
17 keep_prob = tf.placeholder(tf.float32) # dropout (keep probability)
18
19 # Create some wrappers for simplicity
20 def conv2d(x, W, b, strides=1):
21     # Conv2D wrapper, with bias and relu activation
22     x = tf.nn.conv2d(x, W, strides=[1, strides, strides, 1], padding='SAME')
23     x = tf.nn.bias_add(x, b)
24     return tf.nn.relu(x)
25
26
27 def maxpool2d(x, k=2):
28     # MaxPool2D wrapper
29     return tf.nn.max_pool(x, ksize=[1, k, k, 1], strides=[1, k, k, 1],
30                           padding='SAME')
31
32
33 # create Lenet model
34 def Lenet(x, weights, biases, dropout):
35     #reshaping
36     x=tf.reshape(x,shape=[-1,patch_size,patch_size,3])
37     #convolution layer 1
38     conv1=conv2d(x,weights['wc1'],biases['bc1'])
39     # Max Pooling (down-sampling)
40     conv1 = maxpool2d(conv1, k=2)
41
42     # Convolution Layer
43     conv2 = conv2d(conv1, weights['wc2'], biases['bc2'])
44     # Max Pooling (down-sampling)
45     conv2 = maxpool2d(conv2, k=2)
46
47     # Fully connected layer
48     # Reshape conv2 output to fit fully connected layer input
49     fc1 = tf.reshape(conv2, [-1, weights['wd1'].get_shape().as_list()[0]])
50     #fc1=flatten(conv2)
51     fc1 = tf.add(tf.matmul(fc1, weights['wd1']), biases['bd1'])
52     fc1 = tf.nn.relu(fc1)
53     # Apply Dropout
54     fc1 = tf.nn.dropout(fc1, dropout)
55
56     fc2= tf.add(tf.matmul(fc1, weights['wd2']),biases['bd2'])
57     fc2=tf.nn.relu(fc2)
58
59     fc3 = tf.add(tf.matmul(fc2, weights['wd3']), biases['bd3'])
60     fc3 = tf.nn.relu(fc3)
61     # Output, class prediction
62     out = tf.add(tf.matmul(fc3, weights['out']), biases['out'])
63     return out
64
65 """
66 Input
67 x->contains the images in (patch_size,patch_size,3) format
68 The LeNet architecture accepts a 32x32xC image as input, where C is the number
69 of color channels.
70 Modifying to accept 64x64 images
71
72 Architecture
73 Layer 1: Convolutional. The output shape should be 28x28x6.
74
75 Activation. Your choice of activation function.
76
77 Pooling. The output shape should be 14x14x6.
78
79 Layer 2: Convolutional. The output shape should be 10x10x16.

```

```

79     Activation. Your choice of activation function.
80
81     Pooling. The output shape should be 5x5x16.
82
83     Flatten. Flatten the output shape of the final pooling layer such that it's 1D
84     instead of 3D. The easiest way to do is by using tf.contrib.layers.flatten,
85     which is already imported for you.
86
87     Layer 3: Fully Connected. This should have 120 outputs.
88     Activation. Your choice of activation function.
89
90     Layer 4: Fully Connected. This should have 84 outputs.
91     Activation. Your choice of activation function.
92
93     Layer 5: Fully Connected (Logits). This should have 10 outputs.
94     """
95
96
97 #Lenet
98 # Store layers weight & bias
99 lenet_weights = {
100     # 5x5 conv, 1 input, 32 outputs
101     'wc1': tf.Variable(tf.random_normal([5, 5, 3, 6])),
102     # 5x5 conv, 6 inputs, 16 outputs
103     'wc2': tf.Variable(tf.random_normal([5, 5, 6, 16])),
104     # fully connected, 16*16*16 inputs, 400 outputs
105     'wd1': tf.Variable(tf.random_normal([(16)*(16)*16, 400])),
106     # fully connected 400 inputs, 120 outputs
107     'wd2': tf.Variable(tf.random_normal([400,120])),
108     #fully connected 120 inputs, 84 outputs
109     'wd3': tf.Variable(tf.random_normal([120,84])),
110     #fully connected 84 inputs, 2 outputs
111     'out': tf.Variable(tf.random_normal([84, num_classes]))
112 }
113
114
115 lenet_biases = {
116     'bc1': tf.Variable(tf.random_normal([6])),
117     'bc2': tf.Variable(tf.random_normal([16])),
118     'bd1': tf.Variable(tf.random_normal([400])),
119     'bd2': tf.Variable(tf.random_normal([120])),
120     'bd3': tf.Variable(tf.random_normal([84])),
121     'out': tf.Variable(tf.random_normal([num_classes]))
122 }

```

The following figures show the loss and accuracy curves for Lenet. Tensorboard was used here as well to generate the graphs in realtime and was instrumental in hyperparameter tuning.

We can see from the figure that the accuracy increases with time and the loss also decreases, the curves are smooth and convergence is quick. The training time is also faster than the MLP, due to the reduced number of parameters.

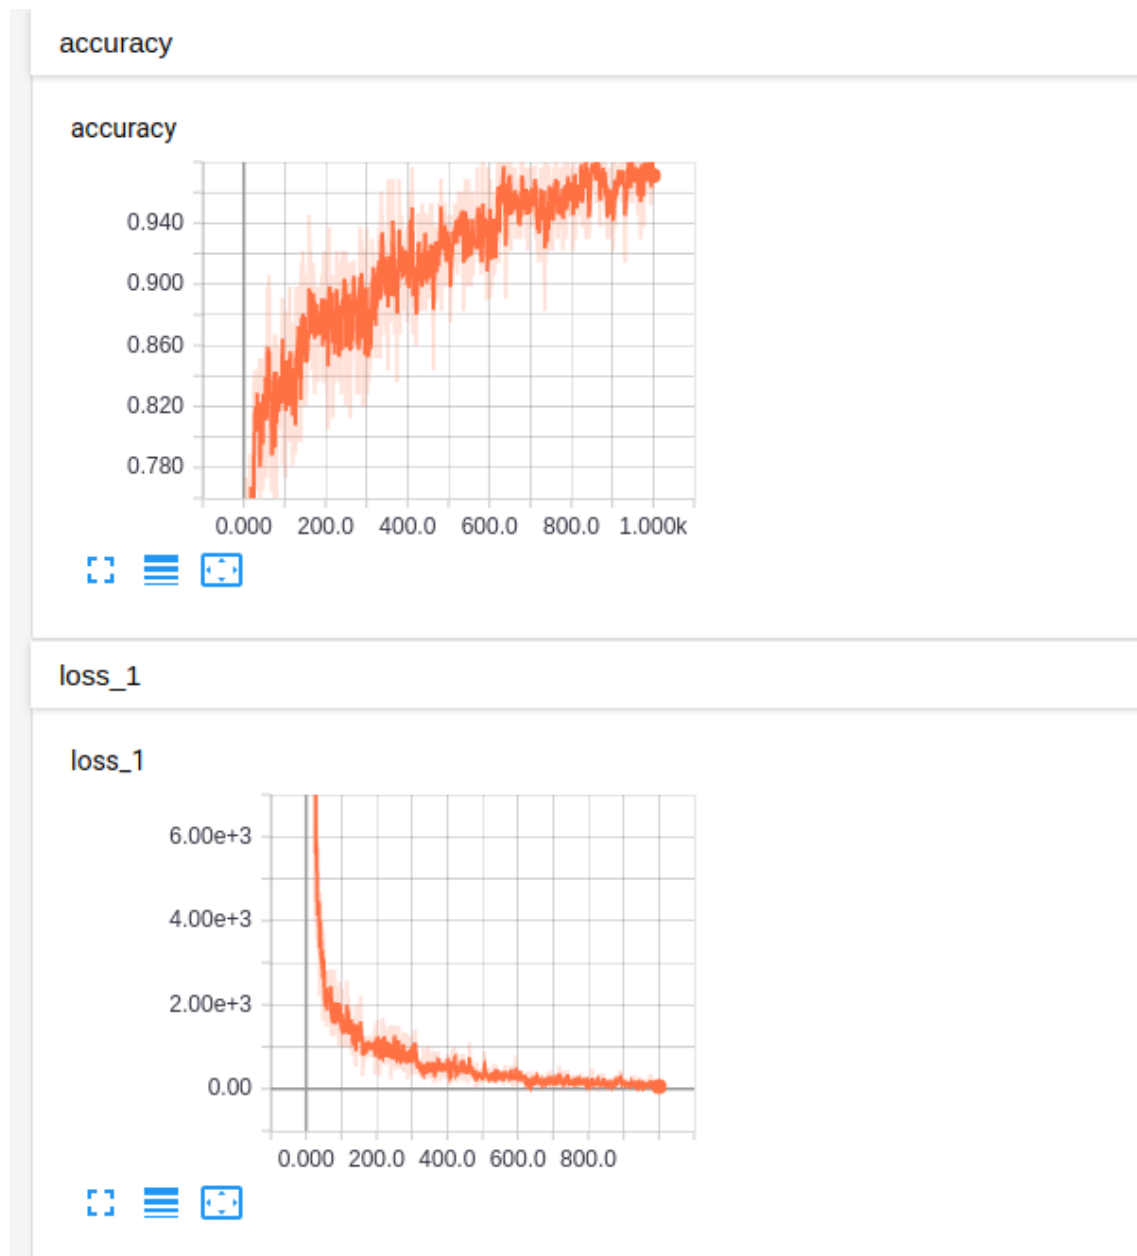


Figure 3: Lenet with batch size 128, 1000 epochs

2.1.3 Own Architecture

The declaration is shown in the code section below.

```

1 }
2
3
4 # Create model
5 def conv_net(x, weights, biases, dropout):
6     # Reshape to match picture format [Height x Width x Channel]
7     # Tensor input become 4-D: [Batch Size, Height, Width, Channel]
8     x = tf.reshape(x, shape=[-1, patch_size, patch_size, 3])
9
10    # Convolution Layer
11    conv1 = conv2d(x, weights['wcl'], biases['bcl'])
12    # Max Pooling (down-sampling)
13    conv1 = maxpool2d(conv1, k=2)
14

```

```

15 # Convolution Layer
16 conv2 = conv2d(conv1, weights['wc2'], biases['bc2'])
17 # Max Pooling (down-sampling)
18 conv2 = maxpool2d(conv2, k=2)
19
20 # Fully connected layer
21 # Reshape conv2 output to fit fully connected layer input
22 fc1 = tf.reshape(conv2, [-1, weights['wd1'].get_shape().as_list()[0]])
23 fc1 = tf.add(tf.matmul(fc1, weights['wd1']), biases['bd1'])
24 fc1 = tf.nn.relu(fc1)
25 # Apply Dropout
26 fc1 = tf.nn.dropout(fc1, dropout)
27
28 # Output, class prediction
29 out = tf.add(tf.matmul(fc1, weights['out']), biases['out'])
30 return out

```

The following figures show the loss and accuracy curves for the custom architecture. Tensorboard was used here as well to generate the graphs in realtime and was instrumental in hyperparameter tuning.

We can see from the figure that the accuracy increases with time and the loss also decreases, the curves are smooth and convergence is quick. The training time is higher than Lenet as the number of parameters are higher. The depth of the convolution layer is higher here, 32 compared to 6 in Lenet. This is for more detailed feature extraction.

As expected, this net produces the best results.

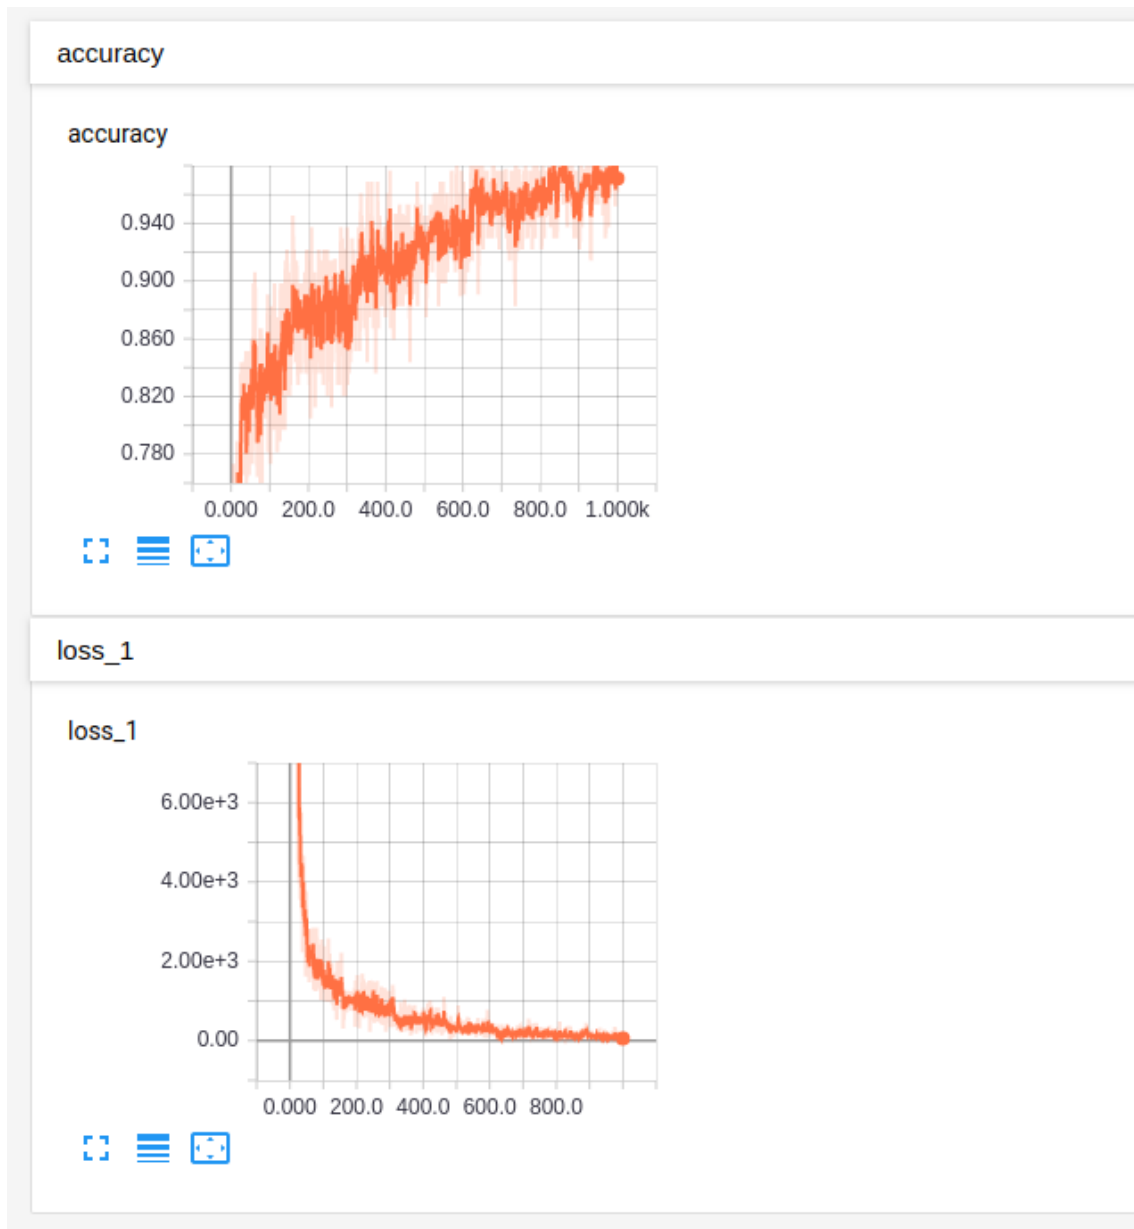


Figure 4: Custom CNN with batch size 128.

3 Conclusion

Overall we can see that the custom CNN architecture performed really well on the face detection task. This was a good primer for the final project of the course.

4 GitHub Repo for Code

The code can be found at <https://github.com/hari0920/ECE-763-Final-Project>.